

Bridging Modalities: Enhancing Cross-Modality Hate Speech Detection with Few-Shot In-Context Learning

Anonymous ACL submission

Abstract

The widespread presence of hate speech on the internet, including formats such as text-based tweets and multimodal memes, poses a significant challenge to digital platform safety. Recent research has developed detection models tailored to specific modalities; however, there is a notable gap in transferring detection capabilities across different formats. This study conducts extensive experiments using few-shot in-context learning with large language models to explore the transferability of hate speech detection between modalities. Our findings demonstrate that text-based hate speech examples can significantly enhance the classification accuracy of vision-language hate speech. Moreover, text-based demonstrations outperform vision-language demonstrations in few-shot learning settings. These results highlight the effectiveness of cross-modality knowledge transfer and offer valuable insights for improving hate speech detection systems.

1 Introduction

Motivation. Hate speech in digital spaces appears in various forms, including text-based tweets and vision-language memes. Researchers have developed models targeting specific modalities (Cao et al., 2023), and recent advancements use large language and multimodal models for more efficient classification (Lin et al., 2024). However, these approaches are often limited to within-distribution datasets and fail to address zero-shot or few-shot scenarios, neglecting the shared characteristics of hate speech across modalities (Hee et al., 2024).

The rise of multimodal hate speech, combining text and visual elements, introduces new challenges (Bates et al., 2023). Existing solutions face two key issues: a scarcity of multimodal datasets due to their recent emergence (Aggarwal et al., 2023) and difficulties in handling copyright issues associated with visual content. Consequently, limited multimodal data adversely affects zero-shot

performance, whereas text-based data, being more abundant and diverse, offers better coverage.

Research Objectives. This paper investigates whether text-based hate speech detection capabilities can be transferred to multimodal formats. By leveraging the richness of text-based data, we aim to enhance the detection of vision-language hate speech, addressing current research limitations and improving performance in low-resource settings.

Contributions. This study makes the following key contributions: (i) We conduct extensive experiments evaluating the transferability of text-based hate speech detection to vision-language formats using few-shot in-context learning with large language models. (ii) We demonstrate that text-based hate speech examples significantly improve the classification accuracy of vision-language hate speech. (iii) We show that text-based demonstrations in few-shot learning contexts outperform vision-language demonstrations, highlighting the potential for cross-modality knowledge transfer. These contributions address critical gaps in existing research and provide a foundation for developing robust hate speech detection systems.

2 Research Questions

As all forms of hate speech share one definition, this study investigates the usefulness of using hate speech from one form, such as text-based hate speech, to classify hate speech in another form, such as vision-language hate speech. Working towards this goal, we formulate two research questions to guide our investigation.

RQ1: Does text hate speech support set help with vision-language hate speech? Visual-language hate speech presents a distinct challenge compared to text-based hate speech, as malicious messages can hide within visual elements or interactions between modalities. It remains uncertain

Dataset	Support		Test	
	# H	# Non-H	# H	# Non-H
Latent Hatred	8189	13,921	-	-
FHM-FG	3,007	5,493	246	254
MAMI	-	-	500	500

Table 1: Statistical distributions of datasets, where "H" represents Hate and "Non-H" represents non-hate

whether text-based hate speech can be useful for classifying visual-language hate speech. We investigate this uncertainty by performing few-shot in-context learning on large language models. This method allows the model to learn from text-based hate speech demonstration examples before classifying visual-language hate speech instances.

RQ2: How does text hate speech support set fare against vision-language hate speech support set? Intuitively, using vision-language hate speech demonstrations should result in superior performance. However, the effectiveness of text-based hate speech demonstrations compared to vision-language hate speech demonstrations remains an open question. To investigate this gap, we conducted another round of few-shot in-context learning on large language models with a vision-language hate speech support set.

3 Experiments

3.1 Experiment Settings

Models. We use Mistral-7B-Instruct-v0.3 (Jiang et al., 2023) and Qwen2-7B-Instruct (Bai et al., 2023) models, which demonstrate strong performance across various benchmarks. Notably, their predecessor models achieve high ELO scores on LMSYS’s Chatbot Arena Leaderboard (Chiang et al., 2024). To facilitate reproducibility and minimize randomness, we use greedy decoding strategy for text generation.

Test Datasets. The Facebook Hateful Memes (FHM) dataset (Mathias et al., 2021) contains synthetic memes categorized into five types of hate incitement: gender, racial, religious, nationality, and disability-based. For evaluation, we use the FHM’s dev_seen split, which includes 246 hateful and 254 non-hateful memes. The Multimedia Automatic Misogyny Identification (MAMI) (Fersini et al., 2022) dataset comprises real-world misogynistic memes classified into shaming, stereotype, objectification, and violence categories. For evaluation, we use the MAMI’s test split, consisting of

500 hateful and 500 non-hateful memes.

Text Support Set. We use the Latent Hatred (ElSherief et al., 2021) dataset, which includes both explicit and implicit forms of hate speech, such as coded and indirect derogatory attacks. This dataset comprises 13,921 non-hateful instances, 1,089 explicit hate speech instances, and 7,100 implicit hate speech instances.

Vision-Language Support Set. We use the FHM dev_seen split for evaluation, containing 3,007 hateful memes and 5,493 non-hateful memes.

3.2 Data Preprocessing

Image Captioning. To perform hateful meme classification with the large language models, we perform image captioning on the meme using the OFA (Wang et al., 2022) model pre-trained on the MSCOCO (Lin et al., 2014) dataset.

Rationale Generation. We employ a common approach to generate informative explanations, directing models to explain their classification decisions (i.e., query -> answer -> explanation). We use Mistral-7B-Instruct-v0.3 to generate detailed rationales for posts in the Latent Hatred, FHM, and MAMI datasets. For the Latent Hatred dataset, we utilize post information and labels, while for the FHM and MAMI datasets, we utilize meme text, captions, and labels. To mitigate noise from varying rationale formulations, we instruct the model to consider both textual and visual elements, focusing on target groups, imagery, and the tweet/meme’s bias perpetuation impact. More details can be found in Appendix F.

3.3 RQ1: Does text hate speech help with vision-language hate speech?

To evaluate the effectiveness of the few-shot in-context learning approach and the Latent Hatred support set, we employed three sampling strategies: Random sampling, TF-IDF sampling, and BM-25 sampling. The TF-IDF and BM-25 strategies leverage the text and caption information of the test record to identify similar examples from the support set, focusing on either the text or the generated rationale. Table 2 shows the comparison of zero-shot and few-shot in-context learning experiment results with Latent Hatred support set.

The experimental results demonstrate that employing a few-shot in-context learning approach with text-based hate speech demonstrations is

Model	# Shots	Dem. Samp.	Matching	FHM			MAMI		
				Acc.	F1	# Invalids	Acc.	F1	# Invalids
Mistral	0-shot	-	-	0.614	0.594	0	0.619	0.568	0
	4-shots	Random	-	0.618	0.613	0	0.655	0.636	0
		TF-IDF	Text.	0.634	0.634	0	0.653	0.649	0
		TF-IDF	Rationale	0.618	0.618	0	0.662	0.658	0
		BM-25	Text.	0.658	<u>0.657</u>	0	0.665	0.662	0
		BM-25	Rationale	0.598	0.596	0	<u>0.676</u>	<u>0.671</u>	0
	8-shots	Random	-	0.620	0.611	0	0.634	0.602	0
		TF-IDF	Text.	0.642	0.641	0	0.665	0.658	0
		TF-IDF	Rationale	0.626	0.625	0	0.657	0.649	0
		BM-25	Text.	0.660	0.658	0	<u>0.685</u>	<u>0.680</u>	0
		BM-25	Rationale	0.612	0.608	0	0.669	0.661	0
	16-shots	Random	-	0.618	0.610	0	0.642	0.611	0
		TF-IDF	Text.	<u>0.644</u>	<u>0.644</u>	0	0.675	0.668	0
		TF-IDF	Rationale	0.632	0.631	0	0.632	0.631	0
		BM-25	Text.	0.638	0.636	0	0.705	0.701	0
		BM-25	Rationale	0.614	0.611	0	0.665	0.659	0
Gwen2	0-shot	-	-	0.624	0.609	0	0.614	0.574	0
	4-shots	Random	-	0.620	0.614	0	0.653	0.632	0
		TF-IDF	Text.	0.632	0.631	0	0.650	0.641	0
		TF-IDF	Rationale	0.634	0.633	0	0.663	0.653	0
		BM-25	Text.	<u>0.644</u>	<u>0.642</u>	0	<u>0.672</u>	<u>0.664</u>	0
		BM-25	Rationale	0.590	0.587	0	0.663	0.654	0
	8-shots	Random	-	0.632	0.628	0	0.645	0.622	0
		TF-IDF	Text.	0.632	0.632	0	0.656	0.650	0
		TF-IDF	Rationale	0.618	0.617	0	0.664	0.656	0
		BM-25	Text.	0.654	0.653	0	0.679	0.674	0
		BM-25	Rationale	0.604	0.603	0	0.654	0.646	0
	16-shots	Random	-	0.632	0.626	0	0.652	0.631	0
		TF-IDF	Text.	0.628	0.628	0	0.656	0.651	0
		TF-IDF	Rationale	0.632	<u>0.631</u>	0	0.665	0.659	0
		BM-25	Text.	0.624	0.624	0	0.678	0.674	0
		BM-25	Rationale	0.630	0.629	0	0.679	0.674	0

Table 2: Comparison of zero-shot and few-shot in-context learning experiment results with *Latent Hatred* support set across different demonstration sampling (Dem. Sampl.) strategies. Underlined results represent the best results within a dataset for the given model and given few-shot setting, while bold results indicate the best results within a dataset for a given model across all few-shot settings. Results highlighted in red, denote few-shot in-context learning results below zero-shot performance.

highly effective in classifying vision-language hate speech. Firstly, while the random sampling strategy could retrieve more irrelevant demonstrations compared to other strategies, the few-shot in-context learning with random sampling surpasses the zero-shot inference performance on both models across two datasets in terms of F1 score. Secondly, the TF-IDF and BM-25 sampling strategies exceed the zero-shot inference performance on both models within the MAMI dataset. Conversely, within the FHM dataset, we observed several instances where some sampling strategies in the few-shot in-context learning scenario performed worse than zero-shot inference. However, these sampling strategies consistently outperformed zero-shot inference when run with 16-shots in-context learning.

Lastly, the best few-shot in-context learning performance within each dataset and each model shows significant improvement over zero-shot model performance. For example, the Mistral model achieves an F1 score improvement of 0.64 and 1.23 on the FHM and MAMI datasets respectively.

3.4 RQ2: How does text hate speech support set fare against vision-language hate speech support set?

Table 3 shows the comparison of zero-shot and few-shot in-context learning experiment results with the FHM support set. The experimental results indicate that using the FHM support set can enhance model performance in some scenarios. However, it is noteworthy that in many instances, few-shot

Model	# Shots	Dem. Samp.	Matching	FHM			MAMI		
				Acc.	F1	# Invalids	Acc.	F1	# Invalids
Mistral	0-shot	-	-	0.614	0.594	0	0.619	0.568	0
	4-shots	Random	-	<u>0.622</u>	<u>0.617</u>	0	0.656	0.642	0
		TF-IDF	Text + Cap.	0.604	0.598	0	<u>0.678</u>	<u>0.670</u>	0
		TF-IDF	Rationale	0.618	0.613	0	0.662	0.652	0
		BM-25	Text + Cap.	0.592	0.584	0	0.662	0.653	0
		BM-25	Rationale	0.620	0.617	0	0.667	0.659	0
	8-shots	Random	-	0.624	0.615	0	0.652	0.632	0
		TF-IDF	Text + Cap.	0.618	0.611	0	0.675	0.664	0
		TF-IDF	Rationale	0.628	0.622	0	0.681	0.670	0
		BM-25	Text + Cap.	0.606	0.599	0	0.672	0.661	0
		BM-25	Rationale	<u>0.628</u>	<u>0.624</u>	0	0.674	0.666	0
	16-shots	Random	-	0.620	0.614	0	0.668	0.651	0
		TF-IDF	Text + Cap.	0.620	0.617	0	0.672	0.665	0
		TF-IDF	Rationale	0.638	0.635	0	0.671	0.661	0
		BM-25	Text + Cap.	0.630	0.625	0	0.682	0.673	0
		BM-25	Rationale	0.634	0.633	0	<u>0.687</u>	<u>0.680</u>	0
Qwen2	0-shot	-	-	0.624	0.609	0	0.614	0.574	0
	4-shots	Random	-	0.606	0.602	0	0.655	0.642	0
		TF-IDF	Text + Cap.	0.620	0.620	0	0.659	0.657	0
		TF-IDF	Rationale	0.636	0.636	0	0.650	0.646	0
		BM-25	Text + Cap.	0.616	0.616	0	0.676	0.674	0
		BM-25	Rationale	0.622	0.622	0	0.669	0.672	0
	8-shots	Random	-	0.592	0.581	0	0.642	0.624	0
		TF-IDF	Text + Cap.	0.606	0.604	0	0.648	0.645	0
		TF-IDF	Rationale	0.620	0.619	0	0.649	0.644	0
		BM-25	Text + Cap.	0.614	0.613	0	0.665	0.662	0
		BM-25	Rationale	<u>0.624</u>	<u>0.623</u>	0	<u>0.669</u>	<u>0.664</u>	0
	16-shots	Random	-	0.602	0.592	0	0.650	0.634	0
		TF-IDF	Text + Cap.	0.610	0.610	0	0.649	0.648	0
		TF-IDF	Rationale	0.604	0.604	0	0.656	0.653	0
		BM-25	Text + Cap.	0.610	0.610	0	0.654	0.653	0
		BM-25	Rationale	<u>0.626</u>	<u>0.626</u>	0	0.653	0.650	0

Table 3: Comparison of zero-shot and few-shot in-context learning experiment results with *FHM* support set across different demonstration sampling (Dem. Sampl.) strategies. Underlined results represent the best results within a dataset for the given model and given few-shot setting, while bold results indicate the best results within a dataset for a given model across all few-shot settings. Results highlighted in red, denote few-shot in-context learning results below zero-shot performance.

in-context learning performs worse than zero-shot model performance when compared against the Latent Hatred support set. Most significantly, the model encounters the most failures on the *FHM* test set despite using the *FHM* train set as a support set. We also observed that the best model performance with the Latent Hatred support set surpasses the best model performance with the *FHM* support set across all instances. We speculate that this discrepancy may stem from the oversimplification of visual information into image captions and the broader topic coverage provided by the extensive Latent Hatred dataset.

4 Conclusion

We investigated the possibility of cross-modality knowledge transfer using few-shot in-context learning with large language models. Our extensive experiment shows that text-based hate speech demonstrations significantly improve the classification accuracy of vision-language hate speech, and using text-based demonstrations in few-shot in-context learning outperforms using vision-language demonstrations. For future works, we aim to extend our analysis to more datasets and explore other cross-modality knowledge transfer approaches such as cross-modality fine-tuning.

Limitations

There are several limitations in this research study

Model Coverage and Model Size. In this study, we evaluated and compared two large models containing 7B parameters. In the future, we aim to extend our analysis to other large models when more computational resources are available.

Large Language Model. In this study, we evaluated few-shot in-context learning in large language models. The experiments are designed in this manner, so to ensure that there can be a fair comparison between the different support sets. We recognize that using a vision-language support set for few-shot in-context learning with a large vision-language model could achieve better performance. However, evaluation using large vision-language models would then be unfair to text support set for few-shot in-context learning.

References

Piush Aggarwal, Pranit Chawla, Mithun Das, Punyajoy Saha, Binny Mathew, Torsten Zesch, and Animesh Mukherjee. 2023. Hateproof: Are hateful meme detection systems really robust? In *Proceedings of the ACM Web Conference 2023*, pages 3734–3743.

Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, Binyuan Hui, Luo Ji, Mei Li, Junyang Lin, Runji Lin, Dayiheng Liu, Gao Liu, Chengqiang Lu, Keming Lu, Jianxin Ma, Rui Men, Xingzhang Ren, Xuancheng Ren, Chuanqi Tan, Sinan Tan, Jianhong Tu, Peng Wang, Shijie Wang, Wei Wang, Sheng-guang Wu, Benfeng Xu, Jin Xu, An Yang, Hao Yang, Jian Yang, Shusheng Yang, Yang Yao, Bowen Yu, Hongyi Yuan, Zheng Yuan, Jianwei Zhang, Xingxuan Zhang, Yichang Zhang, Zhenru Zhang, Chang Zhou, Jingren Zhou, Xiaohuan Zhou, and Tianhang Zhu. 2023. Qwen technical report. *arXiv preprint arXiv:2309.16609*.

Luke Bates, Peter Ebert Christensen, Preslav Nakov, and Iryna Gurevych. 2023. A template is all you meme. *arXiv preprint arXiv:2311.06649*.

Rui Cao, Ming Shan Hee, Adriel Kuek, Wen-Haw Chong, Roy Ka-Wei Lee, and Jing Jiang. 2023. Procap: Leveraging a frozen vision-language model for hateful meme detection. In *ACMMM*.

Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anastasios Nikolas Angelopoulos, Tianle Li, Dacheng Li, Hao Zhang, Banghua Zhu, Michael Jordan, Joseph E. Gonzalez, and Ion Stoica. 2024. *Chatbot arena: An open platform for evaluating llms by human preference*. *Preprint*, arXiv:2403.04132.

Mai ElSherief, Caleb Ziems, David Muchlinski, Vaishnavi Anupindi, Jordyn Seybolt, Munmun De Choudhury, and Diyi Yang. 2021. Latent hatred: A benchmark for understanding implicit hate speech. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 345–363.

Elisabetta Fersini, Francesca Gasparini, Giulia Rizzi, Aurora Saibene, Berta Chulvi, Paolo Rosso, Alyssa Lees, and Jeffrey Sorensen. 2022. Semeval-2022 task 5: Multimedia automatic misogyny identification. In *Proceedings of the 16th International Workshop on Semantic Evaluation (SemEval-2022)*, pages 533–549.

Ming Shan Hee, Shivam Sharma, Rui Cao, Palash Nandi, Preslav Nakov, Tanmoy Chakraborty, and Roy Ka-Wei Lee. 2024. Recent advances in hate speech moderation: Multimodality and the role of large models. *arXiv preprint arXiv:2401.16727*.

Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. 2023. Mistral 7b. *arXiv preprint arXiv:2310.06825*.

Hongzhan Lin, Ziyang Luo, Wei Gao, Jing Ma, Bo Wang, and Ruichao Yang. 2024. Towards explainable harmful meme detection through multimodal debate between large language models. In *Proceedings of the ACM on Web Conference 2024*, pages 2359–2370.

Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. 2014. Microsoft coco: Common objects in context. In *Computer Vision—ECCV 2014: 13th European Conference, Zurich, Switzerland, September 6–12, 2014, Proceedings, Part V 13*, pages 740–755. Springer.

Lambert Mathias, Shaoliang Nie, Aida Mostafazadeh Davani, Douwe Kiela, Vinodkumar Prabhakaran, Bertie Vidgen, and Zeerak Waseem. 2021. Findings of the woah 5 shared task on fine grained hateful memes detection. In *Proceedings of the 5th Workshop on Online Abuse and Harms (WOAH 2021)*, pages 201–206.

Peng Wang, An Yang, Rui Men, Junyang Lin, Shuai Bai, Zhikang Li, Jianxin Ma, Chang Zhou, Jingren Zhou, and Hongxia Yang. 2022. Ofa: Unifying architectures, tasks, and modalities through a simple sequence-to-sequence learning framework. *CoRR*, abs/2202.03052.

A Potential Risks

This project seeks to counteract the dissemination of harmful memes, aiming to protect individuals from prejudice and discrimination based on race, religion and gender. However, we acknowledge the risk of malicious users reverse-engineering memes

to evade detection by CMTL-RAG AI systems, which is strongly discouraged and condemned.

B Licenses and Usage Scientific Artifacts

B.1 Models

All of the LLMs used in this paper contain licenses permissive for academic and/or research use.

- **Mistral-7B-Instruct-v0.3** Apache-2.0 License
- **Qwen2-7B-Instruct** Apache-2.0 License

B.2 Datasets

All of the datasets used in this paper contain licenses permissive for academic and/or research use.

- **Latent Hatred Dataset.** MIT License
- **Hateful Memes Dataset.** MIT License
- **Multimedia Automatic Misogyny Identification.** Creative Commons License (CC BY-NC-SA 4.0)

B.3 Anonymity and Offensive Content

The datasets used in this research contain offensive content, which is crucial for addressing the research questions. Importantly, there are no unique identifiers for the individuals who authored the hateful content in these datasets.

C Computational Experiments

NVIDIA A40 GPUs were utilized for the work done in this paper.

C.1 Experimental Setup

We thoroughly discussed the experimental setup in the main body of the paper. This included descriptions of the models used (Mistral and Gwen2), the number of shots (0-shot, 4-shots, 8-shots, 16-shots), and the different strategies employed for matching (Random, TF-IDF, BM-25) across two datasets (FHM and MAMI). Best-found hyperparameter values were highlighted in the results tables, such as the highest accuracy and F1 scores achieved for each experimental condition.

C.2 Use of Existing Packages

C.2.1 Large Language Models

- **transformers 4.41.1**

C.2.2 Matching and Retrieval Scoring

- **rank-bm25 0.2.2** for BM-25 similarity matching based on this implementation [link](#).
- **scikit-learn 1.5.0** for TF-IDF similarity matching based on this implementation [link](#).

D Few-Shot In-Context Learning

In this approach, we retrieve relevant labelled examples from a 'support dataset' using similarity metrics such as TF-IDF or BM-25 for a given meme from the inference dataset. These examples are then provided as demonstrations in a few-shot prompt to enhance the model's understanding of the meme. Finally, we prompt the model to classify the meme, leveraging the augmented context for improved accuracy.

E Similarity Metrics

E.1 TF-IDF

TFIDF is a statistical measure used to evaluate the importance of a word in a document relative to a collection of documents (corpus). By creating TF-IDF vectors for a 'support dataset', we can use cosine similarity to find the most similar records to a given inference record.

E.2 BM25

BM25 is an advanced version of the TF-IDF weighting scheme used in search engines. It incorporates term frequency saturation and document length normalization to improve retrieval performance. We generate vectors for each record in the 'support dataset' and use cosine similarity to identify records most similar to an inference record.

F Rationale Generation

We use the Mistral-7B-Instruct-v0.3 model, a state-of-the-art language model known for its capabilities in language understanding and generation. We implement a four-shot prompting method to generate explanations for the hateful tweets in the Latent Hatred dataset. Specifically, we select two examples of hateful posts and two non-hateful posts for the four-shot prompt demonstrations. Each demonstration in the prompt follows the following template, given the post text, the post label (hateful or not hateful) and the post rationale:

User: *Determine whether the following post is*

419 *hateful. Text: {text}:*

420 **Assistant:** {label}

421 **User:** *Briefly provide an explanation, in no more*
422 *than three points, for the post being perceived as*
423 *{label}. Your explanation should address the tar-*
424 *geted group, any derogatory imagery or language*
425 *used, and the impact it has on perpetuating bias,*
426 *stereotypes, prejudice, discrimination or inciting*
427 *harm.*

428 **Assistant:** *Answer: {rationale}*

429
430 Similarly, for FHM/MAMI, we select two
431 examples of hateful memes and two non-hateful
432 memes for the four-shot prompt demonstrations.
433 Each demonstration in the prompt follows the
434 following template, given the meme text, the meme
435 label (hateful or not hateful) and rationale:

436
437 **User:** *Determine whether the following meme is*
438 *hateful. Text: {text} Caption: {caption}*

439 **Assistant:** {label}

440 **User:** *Briefly provide an explanation, in no more*
441 *than three points, for the meme being perceived as*
442 *{label}. Your explanation should address the tar-*
443 *geted group, any derogatory imagery or language*
444 *used, and the impact it has on perpetuating bias,*
445 *stereotypes, prejudice, discrimination or inciting*
446 *harm.*

447 **Assistant:** *Answer: {rationale}*

448
449 The demonstration explanation follows a list
450 format, where each list item addresses the targeted
451 group, any derogatory imagery or language
452 used, and the impact it has on perpetuating bias.
453 Finally, to reiterate the classification of the post’s
454 hatefulness, the explanation concludes with the
455 sentence, *"In summary, this post/meme is {label}"*.
456 During inference, the demonstrations provide
457 a structured guide for the model to generate
458 rationales based on the given post/meme’s text and
459 label.